

USCIS sandbox volume test plan — fallback, if no concrete number comes back from USCIS before Monday

Correction, Sep 14 (started running this plan): this doc and an earlier CLOUD_CLAUDE.md note assumed the sandbox mocks any well-formed receipt number. Wrong — the first real run (30 random valid-prefix numbers) came back 100% 404. The sandbox actually serves a small, curated pool of specific receipt numbers. See `src/lib/uscis/sandbox-volume-test.ts`'s `CONFIRMED_SUCCESS_RECEIPTS` (12 confirmed combos) for what actually works — success calls must be drawn from that pool, not generated randomly. Everything else below (run count, pacing, daily total, window) is unchanged and still the live plan.

This is engineering judgment, not a confirmed USCIS requirement. No *sufficiency* number — how much traffic counts as enough — is published anywhere (checked directly: the developer portal, the Torch onboarding PDF, a third-party developer's README, and the Federal Register's PRA notice for the production-access form — none state a volume threshold; see the cloud session's research, folded into "USCIS API access" above). The reply sent to `Torch-Team@uscis.dhs.gov` on Sep 11 asks them directly for a real number. **Check for their answer first, every morning before that day's runs start.** If they give a concrete target, that becomes authoritative — resize the runs below to clearly exceed it, but keep the spread-across-the-day pattern (their own reply said "sufficient traffic (daily)," which implies distributed real usage, not one lump submitted at 9am and never touched again).

Verified ceiling, not a target (Sep 11, cloud session): `developer.uscis.gov/get-started/sandbox` does publish a hard rate-limit ceiling for the Case Status API sandbox — **1,000 requests/day, 5 TPS (1 request per 200ms)**. That's a "don't exceed this" cap, not a "hit this" target — the same page has no language about a required or recommended volume for the 5-day test. **Also checked and ruled out the same day:** a claimed "recommended 100–500 requests/day, 20–50 per batch" figure, sourced to a case-management-software vendor's blog (`legistai.com`), turned out to be fabricated — that page was fetched directly and contains no such numbers and no USCIS citation. Don't treat 100–500/day as a real USCIS figure if it comes up again.

If no answer has come by Monday morning, run this plan as the default.

Why this shape, not just "more calls"

The first attempt (1 success + 1 error call, once a day, 10 calls total across the week) was rejected. Whatever USCIS's internal check actually measures, a single scripted pair executed once a day is trivially distinguishable from a real application generating traffic as real users interact with it throughout the day. **Peter's direction, Sep 11: go for volume near the real ceiling rather than a token increase — 18 runs a day, ~50 transactions per run, spread throughout the day (~900/day total).** Three things ground this, none invented:

1. **~900/day stays under the real 1,000/day ceiling**, with a ~10% margin so an odd retry or a slightly-over-count run doesn't trip a hard lockout. Going all the way to 1,000 exactly would leave no room for error; ~900 is "near the max, less a bit," not "at the max."
2. **Spread across the sandbox's own published operating hours (7AM–8PM ET, Mon–Fri)** — already confirmed real and binding in this project (multiple earlier sessions hit "sandbox down" outside this window; see item 23 and the CW-32/36/38 verification notes). Calls outside this window won't even reach a live sandbox, so all 18 runs have to land inside it — that window is 13 hours, not 16-18; there's no earlier/later time to spread into. 18 runs across 13 hours space out to roughly every 40-45 minutes, 7am to ~8pm ET.
3. **Real variety in receipt numbers and response types**, not the same fixed pair reused every day, and not one flat unpaced burst per run either. This project's own testing already found the sandbox mocks *any* well-formed receipt number (not just the documented `EAC9999103403`), and that different malformed-prefix shapes produce different real error responses (`404` vs `422` , seen across Days 1-4 of the first attempt) — so there's a real, already-proven way to generate legitimately varied traffic.

The plan

Window: Monday, Sep 14 – Friday, Sep 18, 2026. Business days only, matching what USCIS's own reply asked for.

Runs per day: 18, spaced roughly 40-45 minutes apart across the sandbox's actual 7AM–8PM ET operating window (approximately 7:00, 7:45, 8:30am ... through ~7:45pm ET) — not a fixed round-number schedule; stagger the exact start time of each run by a few minutes so it doesn't look like a rigid cron tick every single time.

Per run: ~50 calls, through the app's actual `getCaseStatus()` code path (never a mock/stub), each run internally paced (e.g. roughly 1 call every 1-2 seconds, so a 50-call run takes well under 2 minutes) — comfortably under the 5 TPS ceiling and closer to how a person clicking through case lookups actually generates traffic than a single burst would —

- ~30 success calls, varying the receipt number each time (different valid-format prefixes/numbers, not the same `EAC9999103403` every single call) so responses cover a genuine mix of mock case types and statuses, not one repeated response body.
- ~20 error calls, rotating through different malformed shapes (invalid prefix, wrong length, wrong format entirely) so the error responses are genuinely varied too, not the same `ZZZ999999999` every time.

Daily total: ~900 calls (18 runs × 50 calls) — under the real 1,000/day ceiling with margin, and roughly 90x the 10/week that was rejected. Not claimed as "the number USCIS wants" — still nobody's published sufficiency figure — but it's about as much real, varied, spread-out traffic as this sandbox account can generate in a day without risking the hard quota.

Logging: same discipline as the first attempt — real response codes and bodies, not assumed. Given the volume, a daily rollup per run (call count, success/error split, any real errors or `429` /rate-limit responses hit, and confirmation the day's running total stayed under 1,000) is enough detail for `CLOUD_CLAUDE.md`; no need to transcribe all ~900 calls verbatim the way the 2-call days were. **Watch specifically for any `429` or quota-exceeded response** — that would mean the real ceiling is lower in practice than the published 1,000, and the next day's run count should drop below whatever actually triggered it.

Before re-requesting the affidavit: confirm all 5 days completed with no gaps, at this volume (or higher, if USCIS's reply landed by then with a real number this doesn't already clear). Don't re-send until this is genuinely true; a second rejection costs real credibility with USCIS's team.

What this explicitly is not

Not a guess dressed up as a requirement, and not evidence that more volume is what USCIS actually wants — there's still no published sufficiency threshold, only the 1,000/day ceiling this plan stays under. If Peter or Claude Code finds any better information before Monday — USCIS answers the Sep 11 email, or something else surfaces — that information wins over this plan's numbers. The *shape* (spread across the sandbox's real operating hours, real variety, paced within each run, logged honestly, kept under the real 1,000/day ceiling) is the durable part; the specific run/call counts are Peter's own judgment call (go near the max, less a bit), not a discovered fact.